

Labwork 5: Gaussian Blur Convolution

Nguyen Thanh Binh

October 10, 2025

1 Gaussian blur filter:

To create the filter, we need the size of the filter, which is 7×7 . To get the filter, each pixel is computed by using the normal distribution. μ_x, μ_y are the means of the x and y-axis, which is the mean of the length of the filter.

If even:

$$\mu_x = \frac{H}{2}$$

If odd:

$$\mu_x = \frac{H - 1}{2}$$

For μ_y , replace H with W .

The gaussian filter formula is shown in the following:

$$G(x, y) = \frac{1}{2\pi\sigma^2} \exp \left[-\frac{(x - \mu_x)^2 + (y - \mu_y)^2}{2\sigma^2} \right]$$

Below here is the code implementation:

```
import math
@cuda.jit
def gaussian_blur(layer, stdev):
    tidx = cuda.threadIdx.x
    tidy = cuda.threadIdx.y
    meanx = layer.shape[0] // 2
    meany = layer.shape[1] // 2

    layer[tidx,tidy] = ( 1/( 2*math.pi*(stdev**2))) *
    math.exp(-( (tidx-meanx)**2 + (tidy-meany)**2) / ( 2*(stdev**2) ))
```

`layer` argument is the kernel, and `stdev` is the standard deviation

2 Convolution:

2.1 Without shared memory

The convolution sum is calculated using the formula below, the size of the kernel is 7×7 , then we will have the formula below:

$$\begin{bmatrix} x_{11} & x_{12} & \cdots & x_{17} \\ x_{21} & x_{22} & \cdots & x_{27} \\ \vdots & \vdots & \ddots & \vdots \\ x_{71} & x_{72} & \cdots & x_{77} \end{bmatrix} \begin{bmatrix} y_{11} & y_{12} & \cdots & y_{17} \\ y_{21} & y_{22} & \cdots & y_{27} \\ \vdots & \vdots & \ddots & \vdots \\ y_{71} & y_{72} & \cdots & y_{77} \end{bmatrix}$$
$$= \sum_{i=1}^7 \sum_{j=1}^7 x_{(7-i)(7-j)} y_{(1+i)(1+j)}$$

Below here is the code implementation:

```
@cuda.jit
def convolve(img, layer, res):
    tidx = cuda.threadIdx.x + cuda.blockIdx.x * cuda.blockDim.x
    tidy = cuda.threadIdx.y + cuda.blockIdx.y * cuda.blockDim.y
    # Boundary checking: if the thread coordinates are outside of the image, ignore the thread
    if (i >= tidx) or (j >= tidy):
        return
    delta_tidx = layer.shape[0] // 2
    delta_tidy = layer.shape[1] // 2

    conv_sum = 0
    for i in range (layer.shape[0]):
        for j in range(layer.shape[1]):
            in_x = tidx - i + delta_tidx
            in_y = tidy - j + delta_tidy
            if (in_x >=0) and (in_x < img.shape[0]) and (in_y >=0) and (in_y < img.shape[1]):
                conv_sum += layer[i,j] * img[in_x,in_y]
    res[tidx,tidy] = conv_sum
```

2.2 With shared memory

```
@cuda.jit
def shared_blur(input,output,kernel):
    tile_image = cuda.shared.array(shape = (7,7), dtype = np.uint8)

    tidx = cuda.threadIdx.x + cuda.blockIdx.x * cuda.blockDim.x
    tidy = cuda.threadIdx.y + cuda.blockIdx.y * cuda.blockDim.y
    tile_image[cuda.threadIdx.x, cuda.threadIdx.y] = input[tidx,tidy]
```

```

cuda.syncthreads()

if (tidx >= output.shape[0]) or (tidy >= output.shape[1]):
    return

delta_tidx = kernel.shape[0] // 2
delta_tidy = kernel.shape[1] // 2

conv_sum = 0
for i in range(kernel.shape[0]):
    for j in range(kernel.shape[1]):
        in_x = tidx - i + delta_tidx
        in_y = tidy - j + delta_tidy
        if (in_x >=0) and (in_x < output.shape[0]) and (in_y >=0) and (in_y < output.shape[1]):
            conv_sum += kernel[i,j] * input[in_x,in_y]
output[tidx,tidy] = conv_sum

```

Here are the comparisons before gaussian blurring, and after:



Figure 1: Image before blurring




Figure 2: Image after blurring